

# Enhance Your R Performance and Flexibility with Rcpp

Jon Meek

meekjt (at) gmail.com

meekj (at) ieee.org

Example code on USB sticks and at:

<https://github.com/meekj/RVAR-March2019>

11-March-2019 / RVA R Users

# Overview

- I am not an expert, and other disclaimers...
- Range of native R performance: from horrible to HPC
- “Standard example” is 23 million observations of 5 variables, 813 MB
- Basic benchmarking techniques
- Extending and speeding-up R using C++ with Rcpp
- Today we will only look at single-thread methods
- A real (simple) Rcpp application
- Providing access to a common C library
- Quick mention of other performance issues

# Three Ways to Increment a Vector with Base R: 1

```
> ## Allocate a 23 million point vector
> vlength <- 23e6
> vec <- vector(mode = 'numeric', length = vlength)
> str(vec)
 num [1:23000000] 0 0 0 0 0 0 0 0 0 0 0 ...

> incVal1 <- 1
> ## Use a loop to increment every element
> t_start <- proc.time()
> for (i in 1:length(vec)) {
+   vec[i] <- vec[i] + incVal1
+ }
> proc.time() - t_start
   user  system elapsed 
 1.532   0.046   1.578 
> str(vec)
 num [1:23000000] 1 1 1 1 1 1 1 1 1 1 1 ...
```

## Three Ways to Increment a Vector with Base R: 2 & 3

```
user  system elapsed
1.532   0.046   1.578   for loop from method 1

> ## Do the loop another way
> vec[1:length(vec)] <- vec[1:length(vec)] + incVal1
user  system elapsed
0.306   0.140   0.446

> str(vec)
num [1:23000000] 2 2 2 2 2 2 2 2 2 2 2 ...

> ## Use vectorized R method to increment every element
> vec <- vec + incVal1
user  system elapsed
0.054   0.060   0.114           "The right way"

> str(vec)
num [1:23000000] 3 3 3 3 3 3 3 3 3 3 3 ...
```

# Can we do Better?

- Use Julia for speed? Dirk Eddelbuettel says use Rcpp
- Rcpp provides an easy way to incorporate C++ into R code
- 'for' & 'while' loops in R are slow
  - ▶ vectorize if possible
  - ▶ if not possible use Rcpp
- Other uses for Rcpp
  - ▶ Integrate C/C++ libraries into R for your special requirement
  - ▶ Perform low-level bit-wise calculations
  - ▶ Communicate with hardware (sensors, lab equipment, etc)
  - ▶ Specialized computing where high performance is required
- Try Base R and common packages like dplyr first
- Using R + C++ is similar to how I used FORTRAN + Assembly and Pascal + Assembly in the far past

## Simple Rcpp Code - In-line

```
library(Rcpp)
cppFunction('NumericVector incrementVector(double Increment,
                                           NumericVector TheData) {
    int n = TheData.size(); // C++ way to get length of vector
    for (int i = 0; i < n; ++i) {
        TheData[i] += Increment;
    }
    return TheData;
}')

```

```
> ## Use our simple in-line C++ function to increment every element
> vec <- incrementVector(incVal1, vec)
   user  system elapsed
0.047    0.012    0.058
> str(vec)
 num [1:23000000] 4 4 4 4 4 4 4 4 4 4 4 ...

```

Running this a few times suggests only a minor improvement using C++  
However...

# Do Proper Benchmarking with microbenchmark

- Default is to run code block 100 times (after 2 warm-ups?)
- Result: Classes 'microbenchmark' and 'data.frame'
- Print method provides statistical analysis
- Columns can be added without affecting the print method
- Multiple tests can be combined into a data frame
- \$expr contains the tested expression
- Individual measurements are in \$time
- So, we can make boxplots, etc.
- Also built-in violin plot

# Do Proper Benchmarking - R 3.5.x

Run each example 100 times - Ignore slow methods

```
library(microbenchmark)

> ## Base R - Fast Method
> mb_res1 <- microbenchmark(vec <- vec + incVal1)
> str(vec) ## Note that we got another 100 increments
  num [1:23000000] 104 104 104 104 104 104 104 104 104 104 ...
> ## Rcpp
> mb_res2 <- microbenchmark(vec <- incrementVector(incVal1, vec))
> str(vec)
  num [1:23000000] 204 204 204 204 204 204 204 204 204 204 ...

> ## Look at structure of the microbenchmark result
> str(mb_res1)
Classes microbenchmark and data.frame: 100 obs. of  2 variables:
 $ expr: Factor w/ 1 level "vec <- vec + incVal1": 1 1 1 1 1 1 1 1 1 1 ...
 $ time: num  61967994 61090692 59822301 57745646 57785982 ...

> mb_res <- rbind(mb_res1, mb_res2) # Combine the benchmark results
> mb_res
Unit: milliseconds      expr      min       lq      mean   median      uq      max neval
  vec <- vec + incVal1 57.17667 59.81939 72.34597 60.48852 91.65781 196.0854   100
  vec <- incrementVector(incVal1, vec) 17.87807 18.11191 18.49772 18.31564 18.73056  23.2293   100
```

C++ provides about a 70 % reduction in median run time  
Depending on R instance! (R version, OS version, compiler)



# Do Proper Benchmarking - Nov 2017 - R 3.4.2

Run each example 100 times

```
library(microbenchmark)
```

```
vlength <- 23e6 # Allocate a 23 million point vector  
vec <- vector(mode = 'numeric', length = vlength)
```

```
mb_res1 <- microbenchmark(  
  for (i in 1:length(vec)) {  
    vec[i] <- vec[i] + incVal1  
  }  
)
```

```
mb_res2 <- microbenchmark( vec[1:length(vec)] <- vec[1:length(vec)] + incVal1 )  
mb_res3 <- microbenchmark( vec <- vec + incVal1 )  
mb_res4 <- microbenchmark( vec <- incrementVector(incVal1, vec) )
```

```
rbind(mb_res1, mb_res2, mb_res3, mb_res4)
```


Unit: milliseconds

	expr	min	lq	mean	median	uq	max
for (i in 1:length(vec)) {	vec[i] <- vec[i] + 1	1187.43655	1189.21543	1193.26016	1190.90870	1193.85508	1232.245
vec[1:length(vec)] <-	vec[1:length(vec)] + 1	216.56800	217.42499	221.93871	218.30503	220.16905	338.374
vec <-	vec + 1	25.47441	26.33712	34.28377	26.80986	55.88667	57.498
vec <-	incrementVector(1, vec)	17.27036	17.29239	17.77067	17.50366	18.22541	19.345


(neval = 100 column is cutoff)

C++ provides about a 35 % reduction in median run time

# Incrementing Vector - Base and Rcpp / C++



## microbenchmark Built-in plot



`autoplot(mb_res)`

# A Real Rcpp Application - System or Network Utilization

- Questions often arise well after an “incident”
  - ▶ Why did something slow down, or break?
  - ▶ Too many users or sessions?
  - ▶ Too much bandwidth being consumed?
  - ▶ Was it due to YouTube traffic?
  - ▶ What time of day was the resource stressed? For how long?
- Per session log files typically retained for months
- Packet capture files are too large to retain for long
- Compute estimated throughput or concurrent sessions from network device log files
  - ▶ Millions, or a billion, records
  - ▶ Use session duration and end time
  - ▶ Distribute total bytes, active sessions, or unique users, across one second bins

## Compute Estimated Throughput

~23 million log events covering 24 hours of “end times” (select columns)

```
Time Duration Status BytesSent BytesRecv
2015-04-13T23:57:49 49069 200 401 376
2015-04-13T23:57:49 256 200 522 132
2015-04-13T23:57:49 3063 200 527 3095
2015-04-13T23:57:49 376989 200 398 0
2015-04-13T23:57:49 540 200 766 132
2015-04-13T23:57:49 306792 200 402 0
2015-04-13T23:57:49 802 200 489 196339
...
```

- Use session duration to compute start time
- Distribute bytes received evenly across one second wide bins
- If duration  $\leq 1$  s, full byte count goes in a single bin
- If duration  $> 1$  s, round up to spread across multiple bins
- Two nested loops: Each event; Fill appropriate bins
- R with for loops: 7 - 54 minutes (depending on R version!)
- Rcpp: (as low as) 670 milli-seconds !

## Pre-process Data - Overview

- Reading raw ASCII data with readr is reasonably fast
- Preparing data with Base R & lubridate is very fast
- Simplified data; StartSecond is index / relative time

```
> head(events)
StartSecond Duration BytesRecv
1          2883    49.069         376
2          2932     0.256         132
3          2932     0.269         132
4          2932     0.253         132
5          2929     3.063        3095
6          2556   376.989           0
...
```

- Post-processing with Base R is very fast

# Pure R Code - Quick Look

```
## Read and pre-process data...

## The Loop
for (i in 1:nrow(log_data)) {
  idx <- log_data$StartSecond[i] + 1          # Start index; R starts at 1
  if (log_data$Duration[i] > 1) {             # Does event span multiple bins?
    idt <- as.integer(ceiling(log_data$Duration[i])) # Event duration in bins
    bytes_per_second <- log_data$BytesRecv[i] / idt
    k <- idx + idt - 1                        # Final index to be incremented
    if ((k) > timerange_s) {                  # Don't go past end of vector
      idt <- timerange_s - idx
      k <- idx + idt
    }
    ccu[idx:k] <- ccu[idx:k] + bytes_per_second # Vectorized bin increments
  } else {
    ccu[idx] <- ccu[idx] + log_data$BytesRecv[i] # Single bin to be incremented
  }
}

## Post processing
ccu <- 8 * ccu / 1e3 # 8 bits / byte - kbps

cca_df <- data.frame(Throughput = ccu)          # Make it a dataframe
cca_df$Time <- MinTime + seconds(seq(1:nrow(cca_df)) - 1) # Add time column
```

# Rcpp / C++ Code

## C++ code in it's own file

```
#include <Rcpp.h>
using namespace Rcpp;

// [[Rcpp::export]]
NumericVector concurrentEstimatedThroughput(int outlen, NumericVector StartSecond,
                                             NumericVector Duration, NumericVector Bytes) {

    NumericVector cca(outlen);           // Result will go in this vector
    int k, j, iduration, istart;
    int n = StartSecond.size();          // Number of events
    double bytes_per_second;

    for(int i = 0; i < n; ++i) {          // Process each event
        istart = int(StartSecond[i]);
        iduration = ceil(Duration[i]);    // Number of bins to increment

        if (iduration <= 1) {             // Just increment one bin
            cca[istart] += Bytes[i];
        } else {
            bytes_per_second = Bytes[i] / iduration; // Bytes per bin
            k = istart + iduration - 1;        // Last bin
            if (k >= outlen) {k = outlen - 1;} // Don't go past end of vector
            for (j = istart; j <= k; j++) {    // Distribute bytes across bins
                cca[j] += bytes_per_second;    // covering the event duration
            }
        }
    }
    return cca;
}
```



# Compile and Run C++ Code

```
library(Rcpp)

myPath <- '~/wpl/talks/rvar-201903' # Adjust for local conditions

eventsDataFile <- paste0(myPath, '/events.rds')
codeFile        <- paste0(myPath, '/concurrent_activity.cpp')

events <- readRDS(eventsDataFile)
str(events)

## Compile C++ code from a file
## rebuild & showOutput are optional and mostly useful when messing with compilers and optimization flags

sourceCpp(rebuild = TRUE, showOutput = TRUE, file=codeFile)

## Number of one second wide bins
##
timerange_s <- max(events$StartSecond) - min(events$StartSecond) + 1

## Run the C++ code with timing
##
t_end_prep <- proc.time()
cca <- concurrentEstimatedThroughput(timerange_s, events$StartSecond, events$Duration, events$BytesRecv)
t_end_loop <- proc.time()

t_end_loop - t_end_prep
```

# Post-process & Plot Result

```
library(tidyverse)


str(cca) # Result vector

## Convert to data frame with relative time in seconds and throughput in Mbps
##
cca      <- 8 * cca / 1e6          # 8 bits / byte - Mbps
cca_df   <- data.frame(Throughput = cca) # Make the vector a data frame
cca_df$Time <- seq(1:nrow(cca_df)) - 1 # Add relative time column

str(cca_df)


ggplot(cca_df) +
  geom_point(aes(x = Time, y = Throughput), size = 0.6, color = 'blue', shape = 19) +
  xlab('Time, seconds') + ylab('Throughput, Mbps')
```

## The Result - Estimated Throughput



Estimated throughput for an Internet service over 24 hours with one second granularity.

# MacBook vs “SuperWorkstation”, Estimated Throughput



It's mostly the compiler optimization flags, but gcc is slower for me in 2019.  
(Fall 2017 results shown)

## Selecting the Compiler & Flags for Rcpp

Warning: Can cause problems with package installation

Best to rename when not needed (mv Makevars off-Makevars)

```
cat ~/.R/Makevars
# CC=ccache clang-3.8 -Qunused-arguments
# CXX=ccache clang++-3.8 -Qunused-arguments
# CCACHE_CPP2=yes
# CC=clang-3.8 -Qunused-arguments
CXX=clang++
CXXFLAGS += -Ofast -march=native
# CXXFLAGS += -O3
```

Flags (not compiler) can be set in R:

```
Sys.setenv("PKG_CXXFLAGS"="-Ofast -march=native")
```

Packages have their own Makevars file

## Compute Estimated Throughput - 23 million Events

	Read	Prep	Loop	Post
Base R (for loop)	12.0 s	0.6 s	6.8 - 54 minutes	0.06 s
Python / NumPy	128.6 s	0 s	15.6 minutes	0.33 s
Perl	98.3 s	9.6 s	6.6 minutes	0.30 s
R / C++	12.0 s	0.6 s	0.78 seconds	0.06 s

ASCII data read time includes date / time to seconds conversion

Shorter Base R loop time was for 3.5.2 on Intel NUC Core i7

Python result may not be fair, need to try Pandas

Python read & Prep are done in a single loop

readr is used to read data in R

Perl post time includes writing result

For larger data sets this is an “embarrassingly parallel” computation

# Use Rcpp for Bit Level Computations

- IPAM (IPv4 Address Management)
- `nbPtr <- nbReadAndLoadNetwork(network_description_file)`
- `nbLookupIPaddrs(nbPtr, vector_of_addresses)`
- Finds “shortest” match
- <https://github.com/meekj/netblockr>

Example network description file:

10.16.0.0/12	NOAM xxx North America Supernet
10.16.0.0/22	NOAM PTN Princeton NJ Data Center Servers
10.16.8.0/23	NOAM PTN Princeton NJ West Wing Second Floor
10.18.12.0/23	NOAM SCV Sarah Creek VA
10.48.0.0/12	EMEA xxx EMEA Supernet
10.48.12.0/23	EMEA PSS Portsmouth Southsea
10.48.16.0/24	EMEA ZUR Zurich Wasserschoffi

# Use Rcpp to Access C Library - libpcapR Package

- Set PKG\_LIBS in environment or in Makevars file  
(PKG\_LIBS = -lpcap)
- Add includes to C++ file as usual

```
#include <Rcpp.h>
using namespace Rcpp;
#include <pcap.h>
#include <stdio.h>
#include <string.h>
...
```
- Load network packet capture into a data frame using libpcap
  - ▶ Summarize traffic
  - ▶ Compute throughput with any time granularity
  - ▶ Currently focuses on header data rather than content
  - ▶ Supports IPv4 & IPv6
- <https://github.com/meekj/libpcapR>
- Requires libpcap-dev package to be installed.
- Package needs automated tests, vignette, etc and some users...
- Probably works only on Linux and Mac



# Pre-made Rcpp Packages - Usually Performance Oriented

- dplyr and friends! (transparent use of Rcpp)
- AsioHeaders - Asynchronous network and low-level I/O
- BH - Boost peer-reviewed portable C++ source libraries via headers
- RcppArmadillo - Armadillo Templated Linear Algebra Library
- RcppGSL - GNU Scientific Library
- RcppBDT - Boost Date Time library
- Many, many others

## Rcpp Resources

- I started here: Advanced R Programming by Hadley Wickham: <http://adv-r.had.co.nz/>
- Maybe a better starting point: <http://heather.cs.ucdavis.edu/Rcpp.pdf>
- Full book: Seamless R and C++ Integration with Rcpp by Dirk Eddelbuettel (Springer 2013)
- Rcpp Quick Reference: <https://cran.r-project.org/web/packages/Rcpp/vignettes/Rcpp-quickref.pdf>
- Rcpp Gallery: <http://gallery.rcpp.org/>
- Google → Stackoverflow are your friends, as expected

### General R Performance Resources

- Efficient R Programming, Gillespie and Lovelace, (O'Reilly 2017)
- Performance chapter in Hadley's Advanced R Programming

# C++ Notes

- C++ is a huge language
- Don't need to know a lot of C or C++ to benefit
- Be careful to not index past end of array, etc
- Lots of extensions/updates: C++11, C++14, C++17
- STL has really useful features (expandable containers, etc)
  - ▶ Use the `std::vector<T>` container and the `.push_back(t)` function to grow it
- Boost library
- Free C++ Annotations text (on-line & PDF):  
<http://www.icce.rug.nl/documents/cplusplus/>

# Summary

- Use base R's vectorized functions when possible
- dplyr and other tidyverse packages are fast as well
- Avoid 'for' & 'while' when the loop count is high
- Use a recent version of R and packages
  - ▶ Performance can vary widely between R versions
  - ▶ Newer is not always faster, but sometimes it is much faster!
- Use Rcpp where appropriate
  - ▶ Compiler and flags can make a difference
  - ▶ 4000x performance improvements are possible
- Do benchmarking
- CPU clock speed may suggest how fast R executes base code
- Compiler and flags can have a significant impact on performance
- A busy desktop / laptop will have some effect

# Other Performance Considerations

- Just In Time byte-code compiler enabled by default in R 3.4.0
- Use binary data formats (RDS, FST, Feather, netCDF, etc)
- Read large ASCII flat file(s) once, write single binary file
- Append new ASCII data to existing binary file
- Be sure to save original ASCII data (especially if using fst)
- Hardware can matter, CPU, GPU, etc
- Consider parallelization - R tools are available